

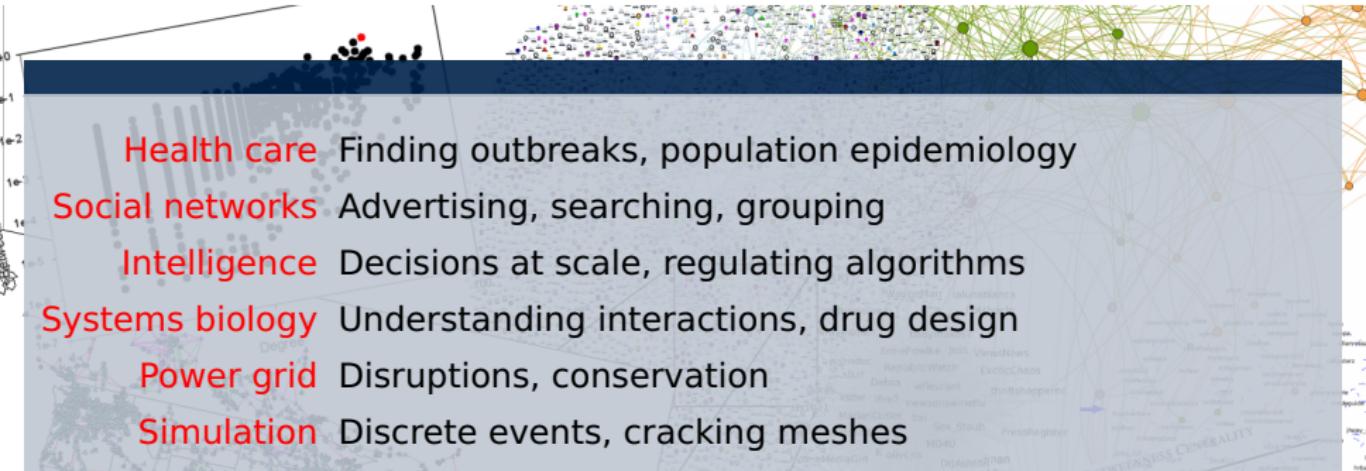


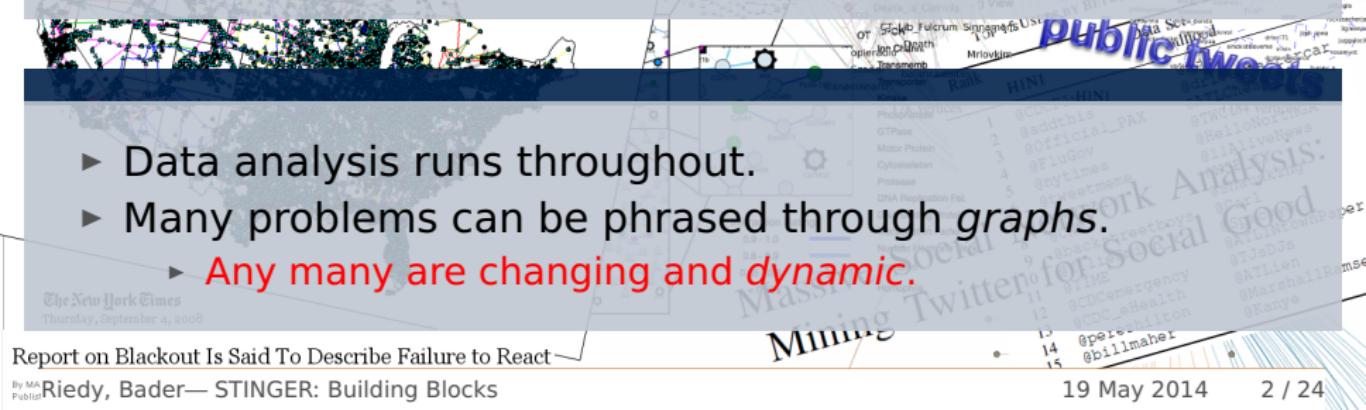
STINGER: Multi-threaded Graph Streaming

Jason Riedy and David Bader

Georgia Institute of Technology

(insert current prefix here)scale data analysis

- 
- Health care** Finding outbreaks, population epidemiology
 - Social networks** Advertising, searching, grouping
 - Intelligence** Decisions at scale, regulating algorithms
 - Systems biology** Understanding interactions, drug design
 - Power grid** Disruptions, conservation
 - Simulation** Discrete events, cracking meshes

- 
- ▶ Data analysis runs throughout.
 - ▶ Many problems can be phrased through graphs.
 - ▶ Any many are changing and dynamic.

The New York Times
Thursday, September 4, 2008

Report on Blackout Is Said To Describe Failure to React

By MA
Riedy, Bader— STINGER: Building Blocks

Mining

19 May 2014

2 / 24

Outline

Motivation: Graph Algorithms for Analysis

Graphs and Streaming Data

STING/STINGER Analysis Framework

Building Blocks for Streaming Graph Data

PageRank

Triangle Counting

Agglomerative Communities

Observations

General approaches

- ▶ High-performance *static graph analysis*
 - ▶ Develop techniques that apply to unchanging massive graphs.
 - ▶ Provides useful after-the-fact information, starting points.
 - ▶ Serves many existing applications well: market research, much bioinformatics, ...
 - ▶ Needs to be $O(|E|)$.
- ▶ High-performance *streaming graph analysis*
 - ▶ Focus on the dynamic changes within massive graphs.
 - ▶ Find trends or new information as they appear.
 - ▶ Serves upcoming applications: fault or threat detection, trend analysis, online prediction...
 - ▶ Can be $O(|\Delta E|)$? $O(\text{Vol}(\Delta V))$? Less data \Rightarrow faster, efficient

Streaming graph data

Data Rates

From www.statisticsbrain.com:

- ▶ 58M posts per day on Twitter (671 / sec)
- ▶ 1M links shared per 20 minutes on Facebook

Other applications (e.g. network security) need to respond nearly at line rate, 81k-1.5M pps on gigabit ethernet.

Opportunities

- ▶ Do not need to analyze the entire graph.
- ▶ Different domains: Throughput & latency
 - ▶ Expose different levels of concurrency
- ▶ Can achieve ridiculous “speed ups.”

Streaming Queries

Different kinds of questions

- ▶ How are individual graph metrics (e.g. clustering coefficients) changing?
- ▶ What are the patterns in the changes?
 - ▶ Are there seasonal variations?
 - ▶ What are responses to events?
- ▶ What are *temporal anomalies* in the graph?
 - ▶ Do key members in clusters / communities change?
 - ▶ Are there indicators of event responses before they are obvious?

On to STING...

Motivation: Graph Algorithms for Analysis

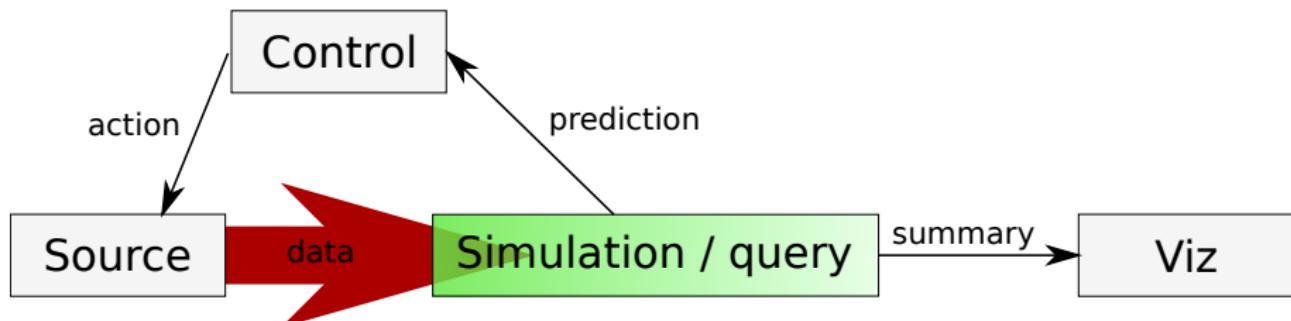
Graphs and Streaming Data

STING/STINGER Analysis Framework

Building Blocks for Streaming Graph Data

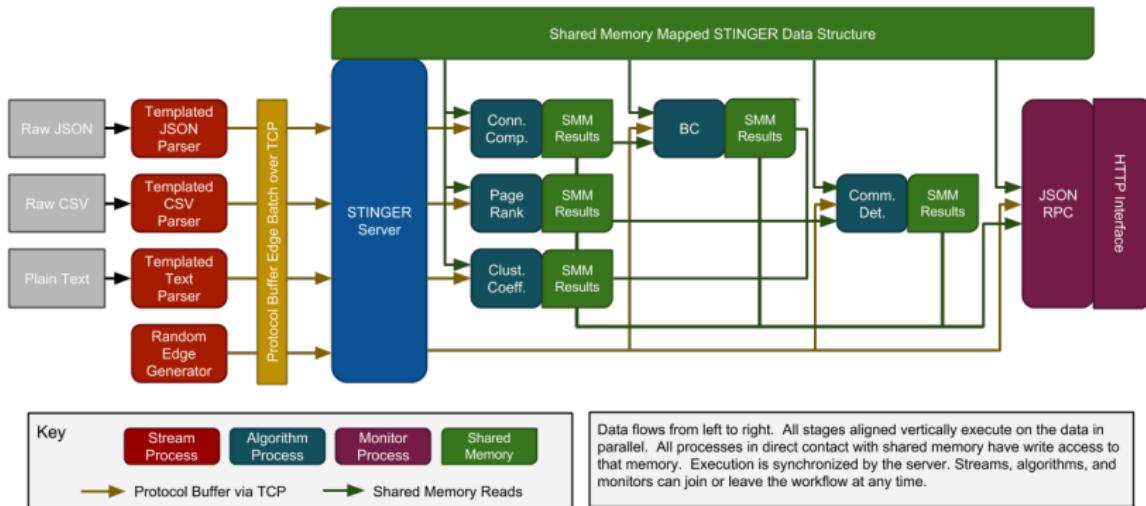
Observations

STING's focus



- ▶ STING: Spatio-Temporal Interaction Networks and Graphs
- ▶ STING manages queries against changing graph data.
 - ▶ Visualization and control often are application specific.
- ▶ Ideal: Maintain many persistent graph analysis kernels.
 - ▶ One current graph snapshot, kernels keep smaller histories.
 - ▶ Also (a harder goal), coordinate the kernels' cooperation.

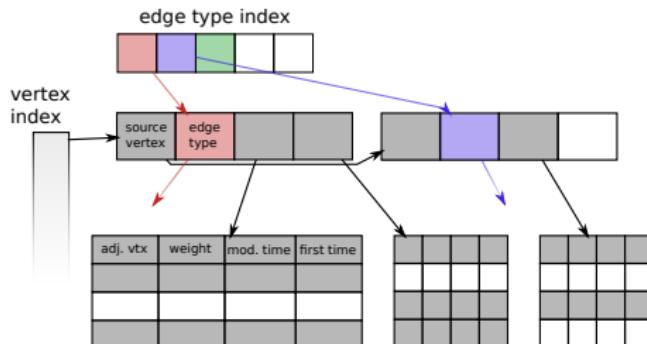
STING: High-level architecture



Slide credit: Rob McColl and David Ediger

- ▶ OpenMP + sufficiently POSIX-ish
- ▶ Multiple processes for resilience

STING Extensible Representation: Core data structure



Initial considerations [Bader, et al.]

- ▶ Be useful for the entire “large graph” community
- ▶ Permit good performance: No single structure is optimal for all.
- ▶ Assume globally addressable memory access and atomic operations
- ▶ Not a *graph database*, but supports types, subsets
- ▶ Large graph \Rightarrow rare conflicts

Building Blocks for Streaming Graph Data

Motivation: Graph Algorithms for Analysis

Graphs and Streaming Data

STING/STINGER Analysis Framework

Building Blocks for Streaming Graph Data

PageRank

Triangle Counting

Agglomerative Communities

Observations

Incremental PageRank

- ▶ PageRank: Well-understood method for ranking vertices based on random walks (related to minimizing conductance).
- ▶ Equivalent problem, solve $(I - \alpha A^T D^{-1})x = (1 - \alpha)v$ given initial weights v .
- ▶ Goal: Use for *seed set expansion*, sparse v .
- ▶ State-of-the-art for updating x when the graph represented by A changes? Re-start iteration with the previous x .
- ▶ Can do significantly better for low-latency needs.
- ▶ Compute the change Δx instead of the entire new x .

Incremental PageRank: Iteration

Iterative solver

Step $k \rightarrow k + 1$:

$$\Delta x^{(k+1)} = \alpha(A + \Delta A)^T(D + \Delta D)^{-1}\Delta x^{(k)} + \alpha[(A + \Delta A)^T(D + \Delta D)^{-1} - A^T D^{-1}]x$$

- ▶ Additive part: Non-zero only at changes.
- ▶ Operator: Pushes changes outward.

Incremental PageRank: Limiting Expansion

Iterative solver

Step $k \rightarrow k + 1$:

$$\Delta \hat{x}^{(k+1)} = \alpha(A + \Delta A)^T(D + \Delta D)^{-1}\Delta \hat{x}_{\text{lex}}^{(k)} + \alpha \Delta \hat{x}_{\text{held}}$$
$$\alpha [(A + \Delta A)^T(D + \Delta D)^{-1} - A^T D^{-1}] \hat{x}$$

- ▶ Additive part: Non-zero only at changes.
- ▶ Operator: Pushes **sufficiently large** changes outward.

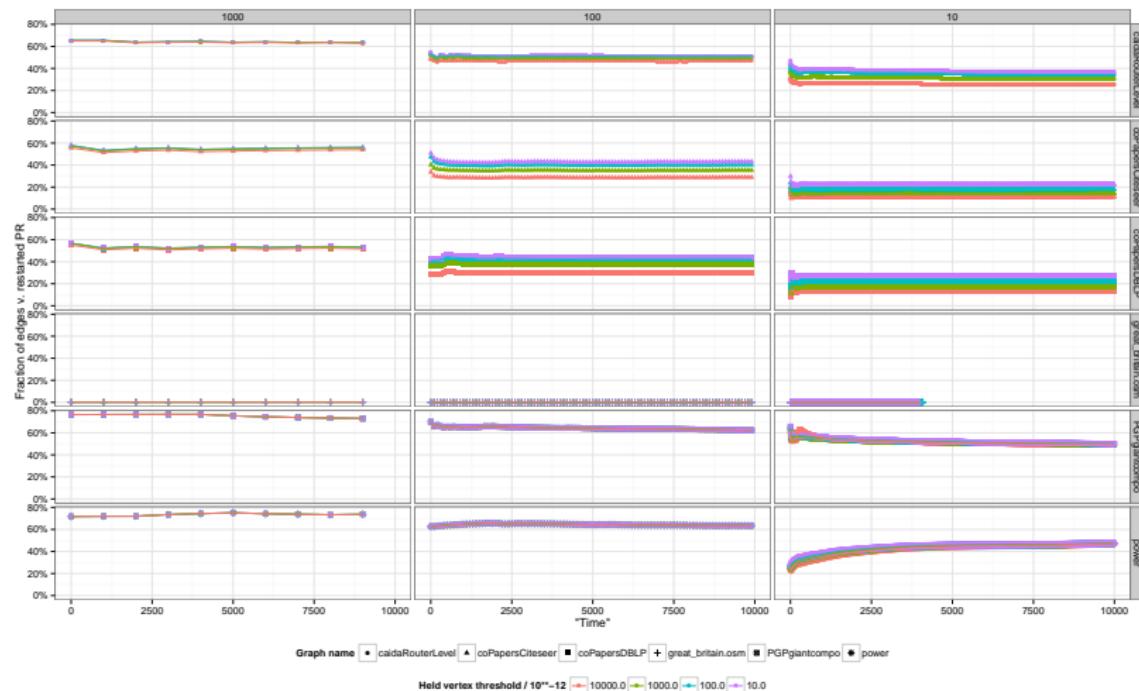
Incremental PageRank: Test Cases

- ▶ Initial, high-level implementation via sparse matrices in Julia.
- ▶ Test graphs from the 10th DIMACS Implementation Challenge.
- ▶ Add uniformly random edges... worst case.
- ▶ Up to 100k in different batch sizes.
- ▶ One sequence of edge actions per graph shared across experiments.
- ▶ Conv. & hold base threshold: 10^{-12}

Graph	V	E	Avg. Deg.	Size (MiB)
caidaRouterLevel	192 244	609 066	3.17	5.38
coPapersCiteseer	434 102	1 603 6720	36.94	124.01
coPapersDBLP	540 486	15 245 729	28.21	118.38
great-britain.osm	7 733 822	8 156 517	1.05	91.73
PGPgiantcompo	10 680	24 316	2.28	0.23
power	4 941	6 594	1.33	0.07

Incremental PageRank: Throughput v. latency

Percent of edge traversals relative to re-started iteration:



Triangle Counting

Current version

- ▶ Count all the triangles around each graph vertex.
- ▶ Used in clustering coefficients (numerator), etc.
 - ▶ Up to 130 000 graph updates per second on X5570 (Nehalem-EP, 2.93GHz)
 - ▶ 2000× speed-up over static recomputation
- ▶ Main algorithm, for each vertex v :
 - ▶ Sort its adjacency list.
 - ▶ For each neighbor w ,
 - ▶ search for w 's neighbors in the sorted list.
- ▶ *Could* compute $\text{diag}(A^3)$, more or less...

Triangle Counting: Small Batches

Low-latency case

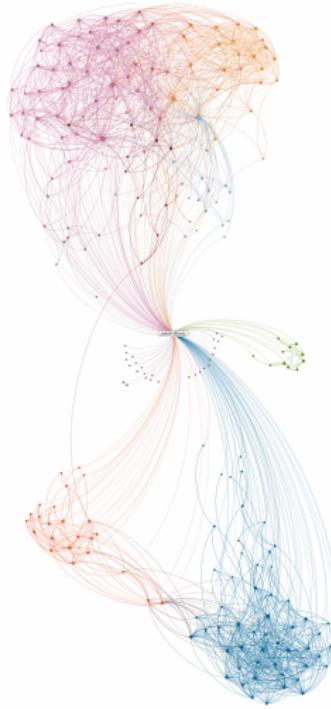
- ▶ In general, $\text{diag}(A^3)$ is a silly option.
- ▶ But $A^3 \Delta x$, a BFS, to count around a few vertices...
- ▶ Brute force (MTAAP10)
 - ▶ Roughly 4x slower with moderate batches, and
 - ▶ less than 2x slower with small batches.
- ▶ Could be reasonable for a quick hack.

Small changes (low latency) may find more applications of linear algebra-like primitives.

Community Detection

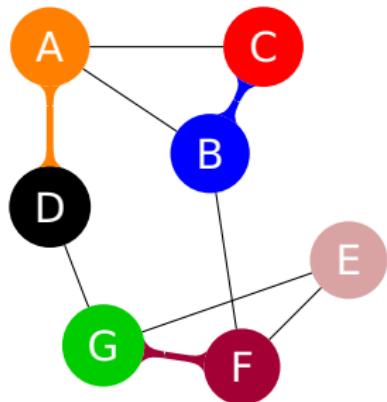
What do we mean?

- ▶ **Partition** a graph's vertices into disjoint communities.
- ▶ Locally optimize some metric, e.g. modularity, conductance
- ▶ Try to capture that vertices are *more similar* within one community than between communities.
- ▶ **Modularity:** More internal edges than expected.



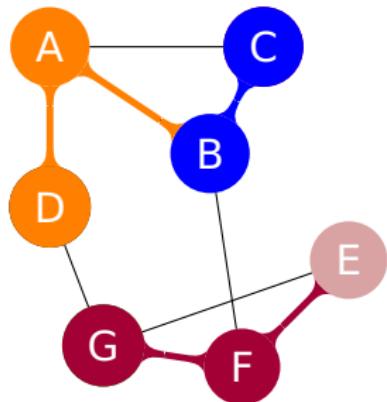
Jason's network via LinkedIn Labs

Parallel Agglomerative Method



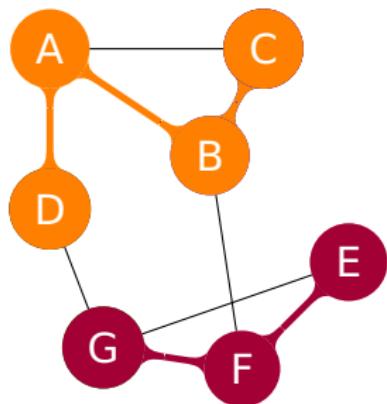
- ▶ Use a **matching** to avoid a serializing queue.
 - ▶ Simple greedy algorithm.
 - ▶ *Would require smuggling data and communication in through element operators...*
- ▶ Highly scalable, $5.8 \times 10^6 - 7 \times 10^6$ edges/sec, $8\times - 16\times$ speed-up on 80-thread E7-8870 (thanks Intel!)
- ▶ Extends to dynamic community maintenance
 - ▶ Extract vertices from communities, re-agglomerate
 - ▶ *Matrix triple product-ish*

Parallel Agglomerative Method



- ▶ Use a **matching** to avoid a serializing queue.
 - ▶ Simple greedy algorithm.
 - ▶ *Would require smuggling data and communication in through element operators...*
- ▶ Highly scalable, $5.8 \times 10^6 - 7 \times 10^6$ edges/sec, $8\times - 16\times$ speed-up on 80-thread E7-8870 (thanks Intel!)
- ▶ Extends to dynamic community maintenance
 - ▶ Extract vertices from communities, re-agglomerate
 - ▶ *Matrix triple product-ish*

Parallel Agglomerative Method



- ▶ Use a **matching** to avoid a serializing queue.
 - ▶ Simple greedy algorithm.
 - ▶ *Would require smuggling data and communication in through element operators...*
- ▶ Highly scalable, $5.8 \times 10^6 - 7 \times 10^6$ edges/sec, $8\times - 16\times$ speed-up on 80-thread E7-8870 (thanks Intel!)
- ▶ Extends to dynamic community maintenance
 - ▶ Extract vertices from communities, re-agglomerate
 - ▶ *Matrix triple product-ish*

Seed Set Expansion

Problem

- ▶ Given a small number of vertices, and
- ▶ find a region of interest around them.

- ▶ Start with a subset consisting of the selection.
- ▶ Evaluate the change in *modularity* around the current subset.
- ▶ Absorb all vertices that...
 - ▶ may increase modularity by a significant amount, or
 - ▶ are within the top 10% of changes, or...
- ▶ Repeat until the set is large enough.
- ▶ *Step-wise guided expansion doesn't fit current primitives.*

Observations

- ▶ Throughput / latency trade offs:
 - ▶ Different levels of parallelism and optimizations
 - ▶ Larger batches \Rightarrow higher throughput, more collisions
 - ▶ Small batches \Rightarrow lower latency, more scattered
 - ▶ Impact optimizations similarly to direction-optimized BFS
- ▶ Can build proposed building blocks against STINGER
- ▶ Many algorithms are not naturally expressed:
 - ▶ Matching
 - ▶ Guided set expansion by changing criteria
 - ▶ Streaming versions of these...
- ▶ Targets for version 2?

STINGER: Where do you get it?



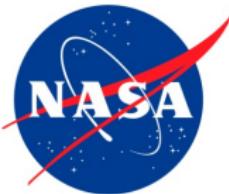
www.cc.gatech.edu/stinger/
Gateway to

- ▶ code,
- ▶ development,
- ▶ documentation,
- ▶ presentations...
- ▶ (working on usage and development screencasts)

Remember: Still academic code, but maturing.

Users / contributors / questioners:
Georgia Tech, PNNL, CMU, Berkeley,
Intel, Cray, NVIDIA, IBM, Federal
Government, Ionic Security, Citi

Acknowledgment of support



NORTHROP GRUMMAN



CRAY



TOSHIBA